

# Advanced Techniques for Combinatorial Algorithms: External-Memory Algorithms

Gianluca Della Vedova

Univ. Milano-Bicocca  
<https://gianluca.dellavedova.org>

April 28, 2020

# Memory Hierarchy

- CPU Registers
- L1 Cache: 32 – 256 KBytes, latency  $< 10^{-9}$  secs
- L2 Cache: 1 – 16 MBytes, block transfer size  $B = 32$  Bytes
- RAM: 1 – 32 GBytes,  $B = 64$  Bytes
- Disk: 100 GBytes - 1 PBytes,  $B = 8$  KBytes, latency  $> 10^{-3}$  secs

# PDM model

- Parallel Disk Model
- Locality of reference
- Parallel disk access
- Disk striping (data across several disks)
- Count I/O operations

# PDM parameters

- $N$  = problem size (in units of data items)
- $M$  = internal memory size (in units of data items)
- $B$  = block transfer size (in units of data items)
- $D$  = number of independent disk drives;
- $P$  = number of CPUs
- $Q$  = number of queries (for a batched problem);
- $Z$  = answer size (in units of data items).
- $M < N$  and  $1 \leq DB \leq M/2$
- $n = N/B$ ,  $m = M/B$ ,  $q = Q/B$ ,  $z = Z/B$

# Basic operations

- Scan:  $\Theta(\frac{N}{DB}) = \Theta(\frac{n}{B})$
- Sort:  $\Theta(\frac{N}{DB} \log_{M/B} \frac{N}{B}) = \Theta(\frac{n}{D} \log_{M/B} n)$
- Search:  $\Theta(\log_{DB} N)$
- Output:  $\Theta(\max\{1, \frac{Z}{DB}\}) = \Theta(\max\{1, \frac{z}{D}\})$

# Disk striping

- I/O only on entire stripes
- cohesive set of disks
- $D$  disks as a logical disk with logical block size  $DB$

## Main idea

- 1 disk: each I/O step transmits one block of size  $DB$
- $D$  disks: each I/O step consists of  $D$  simultaneous block transfers of size  $B$  each.
- Same number of I/O steps

# Distribution sort

## $S$ buckets

- By choosing  $S - 1$  pivots
- needs buckets of similar size, so  $O(\log_S n)$  recursion layers
- scan to build the buckets. When a buffer is full  $\Rightarrow$  write it
- $O(m)$  buckets
- probabilistic approach to select the pivots

# Multiway Partitioning (PDM)

## Multiway Partitioning

- $M = \{m_1, \dots, m_d\}$  ordered set of pivots
- $S$ : unordered set of elements
- $A_i$ :  $i$ -th bucket.  $a_i \in A_i$ ,  $m_{i-1} < a_i \leq m_i$
- Goal: Compute  $A_i$ s
- Goal: Compute  $|A_i|$



# Multiway Partitioning (PDM)

---

## Algorithm 1: MultiPartition

---

```
1 Split  $A$  into sets  $S_1, \dots, S_P$ ;  
2 foreach processor  $i$  in parallel do  
3   Read the vector of pivots  $M$  into the cache;  
4   Partition  $S_i$  into  $d$  buckets,  $J_i =$  number of items in each bucket  
5 Prefix Sums on  $\{J_1, \dots, J_P\}$  in parallel;  
6 foreach processor  $i$  in parallel do  
7   Write elements  $S_i$  into memory locations offset appropriately by  $J_{i-1}$  and  $J_i$   
8 compute  $|A_i|$ s, using the prefix sums stored in  $J_P$ 
```

---

# PDM references

- Jeffrey Scott Vitter. Algorithms and Data Structures for External Memory. Foundations and Trends in Theoretical Computer Science. Now Publishers, 2008 [http://www.ittc.ku.edu/~jsv/Papers/Vit.IO\\_book.pdf](http://www.ittc.ku.edu/~jsv/Papers/Vit.IO_book.pdf)
- L. Arge, M. T. Goodrich, M. Nelson, and N. Sitchinava. Fundamental parallel algorithms for private-cache chip multiprocessors. In Proceedings of SPAA '08, pages 197–206. ACM, 2008.
- Fundamental parallel algorithms for private-cache chip multiprocessors <https://dl.acm.org/citation.cfm?id=1378573>