

Advanced Techniques for Combinatorial Algorithms: Data Streams and Map-Reduce

Gianluca Della Vedova

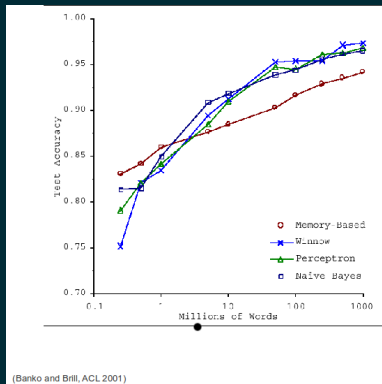
Univ. Milano-Bicocca
<https://gianluca.dellavedova.org>

May 18, 2020

Gianluca Della Vedova

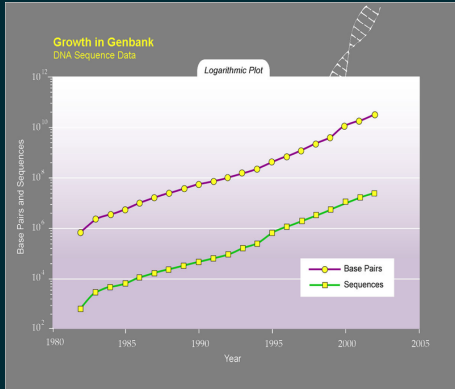
- Advanced Techniques for Combinatorial Algorithms
- <https://gitlab.com/dellavg/advanced-algorithms>
- <https://gianluca.dellavedova.org>
- gianluca.dellavedova@unimib.it

Fact 1




- Huge data are among us

And more are coming

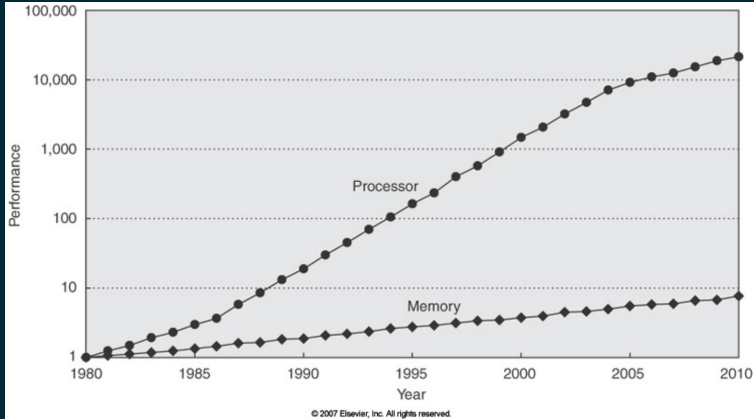


From 1982 to the present, the number of bases in GenBank has doubled approximately every 18 months (<ftp://ftp.ncbi.nih.gov/genbank/gbrel.txt>).

Moore's Law



Fact 2



- Moore's Law is unfair.

Problem

- Input data too large for a single computer
- Don't fit into memory

Solution

- Split data into parts
- If only it were so easy
- Embarassingly parallel problems

Solutions

- parallel algorithms
- map reduce
- data streaming
- External-memory algorithms
- Are all related

Map Reduce model

- **Input:** hash = $\langle \text{key} \mapsto \text{value} \rangle$ pairs
- Map
- Shuffle
- Reduce
- **Output:** hash

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash
- Shuffle: all values with same key k are assigned to a unique processor
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash
- Shuffle: all values with same key k are assigned to a unique processor
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash
- Shuffle: all values with same key k are assigned to a unique processor
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash in parallel
- Shuffle: all values with same key k are assigned to a unique processor
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash in parallel
- Shuffle: all values with same key k are assigned to a unique processor automatic
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs

Map Reduce model

- Mapper: receives a $\langle \text{key} \mapsto \text{value} \rangle$ pair, computes a hash in parallel
- Shuffle: all values with same key k are assigned to a unique processor automatic
- Reducer: receives all values with same value k , computes a multiset of $\langle k, \text{value} \rangle$ pairs sequential algorithm on the values

k -th Frequency Moment

Instance

k , a list $L = \langle x_1, \dots, x_n \rangle$ of elements over alphabet Σ

Output

$\sum_{\sigma \in \Sigma} f^k(\sigma)$, where $f(\sigma)$ is the number of occurrences of *sigma* in L

Example

$L = \langle 0, 1, 1, 1, 2, 0, 1, 2, 0, 1, 1, 2 \rangle$

Output, $k = 2$

$$3^2 + 6^2 + 3^2 = 54$$

Example: k -th frequency moment

Algorithm 1: k -FrequencyMoment

Data: k , a list $= L \langle x_1, \dots, x_n \rangle$

```
1  $\mu_1(\langle i; x_i \rangle) = \langle x_i; i \rangle$  /*  $i$  is the index */
2  $\rho_1(\langle x_i; v_1, \dots, v_m \rangle) = \langle x_i; m^k \rangle$ ; /* All occurrences of a symbol are
   grouped together */
3  $\mu_2(\langle x_i; v \rangle) = \langle \$; v \rangle$ ; /* Time to sum */
4  $\rho_1(\langle \$; v_1, \dots, v_l \rangle) = \langle \$; \sum v_i \rangle$ ;
```

Rationale

- No machine can store whole input
- Too many machines = **bad**
- Shuffling is expensive
- No shared memory
- Loose synchronization

Efficient Algorithm

- at most $O(n^{1-\epsilon})$ machines
- each mapper is a RAM with $O(n^{1-\epsilon})$ space, $\text{poly}(n)$ time
- each reducer is a RAM with $O(n^{1-\epsilon})$ space, $\text{poly}(n)$ time
- polylogarithmic mapreduce rounds
- reducers have $O(n^{2-2\epsilon})$ overall space
- each key has $O(n^{1-\epsilon})$ values

Classes

- \mathcal{MRC}^i if $O(\log^i n)$ rounds
- $\mathcal{MRC} = \bigcup_{i=0}^{\infty} \mathcal{MRC}^i$
- $\mathcal{MRC}^0, \mathcal{MRC}^1$

Map Reduce model

Howard J. Karloff, Siddharth Suri, Sergei Vassilvitskii: A Model of Computation for MapReduce. SODA 2010: 938-948

Relations between models

Map Reduce and PRAM

Reconcile inputs

$$x_1, \dots, x_n \mapsto \langle 1, x_1 \rangle, \dots, \langle n, x_n \rangle$$

Theorem 4.1, Karloff et al, SODA 2010

If $\mathcal{P} \neq \mathcal{NC}$ then $\mathcal{MRC} \not\subseteq \mathcal{NC}$.

Proof

- Pick a \mathcal{P} -complete problem, and pad its input
- Use a single reducer to solve it

Map Reduce and PRAM

Theorem 7.1, Karloff et al, SODA 2010

Any CREW PRAM algorithm using $O(n^{2-2\epsilon})$ total memory, $O(n^{2-2\epsilon})$ processors and $t(n)$ time can be run in $O(t(n))$ rounds in \mathcal{MRC} .

Proof

- processor \rightarrow reducer
- memory location \rightarrow reducer
- shuffle to coordinate read requests

Map Reduce and PDM

Shuffle step

- Matrix: R =rows=reduce steps, M =columns=keys
- sparse matrix
- Layout depending on maps
- Rearrange into a column-wise ordering

Graph Algorithms

- For each edge e , store its successor $s(e)$
- Euler tour of G = edge-disjoint cycles
- merge any two cycles with a common vertex u

Minimum Spanning Tree (MRC)

Algorithm 2: MST

Data: Array

- 1 dense graph $G = \langle V, E \rangle$
/* $|V| = n, |E| \geq n^{1+c}$ */
 - 2 $V_1, \dots, V_k \leftarrow$ random balanced partition of V ;
 - 3 $E_{i,j} \leftarrow \{(u, v) \in E \mid u, v \in V_i \cup V_j\}$;
 - 4 $G_{i,j} \leftarrow$ subgraph of G induced by $E_{i,j}$;
 - 5 **foreach** $G_{i,j}$ **do**
 - 6 $M_{i,j} \leftarrow$ minimum spanning forest of $G_{i,j}$
 - 7 $H \leftarrow \langle V, \bigcup_{i,j} M_{i,j} \rangle$;
 - 8 $M \leftarrow$ minimum spanning tree of H ;
-

Minimum Spanning Tree (MRC)

MST Algorithm is correct

- 1 Let $e \in E \setminus E(H)$.
- 2 Then e is a heaviest edge in some cycles of $G_{i,j}$.
- 3 The same cycle is also in G .

Minimum Spanning Tree (MRC)

$k = n^{c/2}$. Then $|E_{i,j}|$ is $\tilde{O}(n^{1+c/2})$

- 1 $W_i = \{v \in V : 2^{i-1} < \deg(v) \leq 2^i\}$
- 2 If $|W_i| < 2n^{c/2} \log n$ then $\sum_{v \in W_i} \deg(v) = \tilde{O}(n^{1+c/2})$
- 3 Else $|W_i \cap V_j| = O(\log n)$ via Chernoff bound. Proof completed via union bound

Chernoff bound

- 1 $X = \sum_i^q X_i$, X_i ind. random variables, $Pr[X_i = 1] = Pr[X_i = -1] = 1/2$
- 2 $Pr[X \geq a] \leq e^{\frac{-a^2}{2q}}$

MRC General Technique

Def. f parallelizable function

- 1 If there exists functions g, h such that:
- 2 For any partition T_1, \dots, T_k of the universe set S , $f(S) = h(g(T_1), \dots, g(T_k))$
- 3 g, h are polynomial time computable
- 4 Output of g has $O(\log n)$ bits.

MRC algorithm for f_1, \dots, f_k

- 1 S_1, \dots, S_k subsets of S , $k \leq n^{2-3\epsilon}$, $\sum |S_i| \leq n^{2-2\epsilon}$
- 2 $f_1(S_1), \dots, f_k(S_k)$ with $O(n^{1-\epsilon})$ reducers, $O(n^{1-\epsilon})$ space each.

k -th frequency moment

Algorithm 3: k -FrequencyMoment

Data: k , a sequence $\langle x_1, \dots, x_n \rangle$

- 1 $\mu_1(\langle i; x_i \rangle) = \langle x_i; i \rangle;$
 - 2 $\rho_1(\langle x_i; v_i, \dots, v_m \rangle) = \langle x_i; m^k \rangle;$
 - 3 $\mu_2(\langle x_i; v \rangle) = \langle \$; v \rangle;$
 - 4 $\rho_1(\langle \$; v_i, \dots, v_l \rangle) = \langle \$; \sum v_i \rangle;$
-

k -th frequency moment

Problem: all elements can be the same

$O(n)$ space, instead of $O(n^{1-\epsilon})$

Solution

- $g(\{t_1, \dots, t_k\}) = k$
- $h(\{i_1, \dots, i_m\}) = (i_1 + \dots + i_m)^k$

Proof

- $h(g(T_1), \dots, g(T_m)) = |S|^k$
- S set of pairs with same value

Maximal Matching (MRC)

Matching

- Let $G = \langle V, E \rangle$ be a graph.
- Then $M \subseteq E$ is a matching if no two edges of M shares a vertex
- M is maximal if $\nexists M_1 \supseteq M$ s.t. M_1 is a matching

Maximal Matching (MRC)

Algorithm 4: MaximalMatching

Data: Array

```
1 dense graph  $G = \langle V, E \rangle$ 
2  $M \leftarrow \emptyset, S = E$ ;
3  $E' \leftarrow$  random sample of  $S$ ;
4 if  $|E'| > \eta$  then Fail;
5  $M' \leftarrow$  maximal matching on  $E'$  /* 1 reducer */
6  $M \leftarrow M \cup M'$ ;
7 Remove all edges conflicting with  $M$ ;
8  $S \leftarrow$  remaining edges;
9 if  $|S| > \eta$  then go to step 2;
10  $M' \leftarrow$  maximal matching on  $S$  /* 1 reducer */
11  $M \leftarrow M \cup M'$ ;
```

Connected components

Algorithm 5: ConnectedComponents

```
1 Label each vertex either up or down;  
2 foreach edge (v, w) do  
3   |   Add the component of the down vertex to that of the up vertex;  
   |   /* The root of the lower component is now a child of the root  
   |     of the upper component                                     */  
4 foreach vertex v do  
5   |    $\text{parent}[v] \leftarrow \text{parent}[\text{parent}[v]]$ 
```

Connected components

```
def map(u, v, p1, p2):  
    if p1 == p2: #same component  
        return []    #do nothing  
    else:  
        h1 = hash(p1, r)%2  
        h2 = hash(p2, r)%2  
        if h1 != h2:  
            return [(p1,p2) if h1 else (p2,p1)]  
  
def reduce(list):  
    return list[0]
```

Additional Bibliography on Map Reduce

On distributing symmetric streaming computations

<http://portal.acm.org/citation.cfm?doid=1824777.1824786>

Counting triangles and the curse of the last reducer

<http://portal.acm.org/citation.cfm?doid=1963405.1963491>

A Model of Computation for MapReduce <http://dl.acm.org/citation.cfm?id=1873677>

Data Streams

Almost permutation

- Input: a permutation π of $\{1, \dots, n\}$ with one element missing
- Find the missing element
- Streaming: read the input once
- $\lceil \log n \rceil$ bits of memory

Almost permutation

Algorithm

- Keep the parity bits
- and add $\sum_{i=1}^n i$

Reservoir sampling

- Take k elements from a stream
- Add to S the first k elements
- The i -th elements is kept with probability k/i
- In case, remove a random element of S

Uniform sampling

Time t , $i \leq t$. Then $Pr[x_i \in S] = \frac{s}{t}$

Count-Min sketch

- Turnstile model
- Approximate counting: number of occurrences of a symbol
- $d \times w$ matrix *count*[]
- $w = e/\epsilon$, $d = \log 1/\delta$. We want additive error $\leq \epsilon$ with probability $\geq 1 - \delta$
- d hash functions $h_1, \dots, h_d : \{1, \dots, N\} \mapsto \{1, \dots, w\}$
- update (j, l_i) , then $count[k, h_k(j)] \leftarrow count[k, h_k(j)] + l_i, \forall k$
- query $A[i]$, answer $\min_j count[j, h_j(i)]$

HyperLogLog sketch

- Estimate the cardinality of a set of integers
- $p \leftarrow$ Max number of leading zeroes
- Estimate 2^p
- Go through hash function to manage arbitrary sets
- Split into subsets to reduce variance